

CS146: Data Structures and Algorithms

Lecture 9



**LINEAR TIME SORTING
LOWER BOUNDS**

**INSTRUCTOR: KATERINA POTIKA
CS SJSU**

Sorting So Far – 1st Algorithm

2

- Insertion sort:
 - Easy implementation
 - Fast on small inputs (less than ~50 elements)
 - Fast on nearly-sorted inputs
 - $O(n^2)$ worst case
 - $O(n^2)$ average (equally-likely inputs) case
 - $O(n^2)$ reverse-sorted case

Sorting So Far – 2nd Algorithm

3

- Merge sort:
 - Divide-and-conquer:
 - ◆ Split array in half
 - ◆ Recursively sort subarrays
 - ◆ Linear-time merge step
 - $O(n \lg n)$ worst case
 - Doesn't sort in place

Sorting So Far – 3rd Algorithm

4

- Heap sort:
 - Uses the heap data structure
 - ◆ Complete binary tree
 - ◆ Heap property: parent key $>$ children's keys
 - $O(n \lg n)$ worst case
 - Sorts in place
 - Fair amount of shuffling memory around

Sorting So Far – 4th Algorithm

5

- Quick sort:
 - Divide-and-conquer:
 - ◆ Partition array into two subarrays, recursively sort
 - ◆ All elements of first subarray < all elements of second subarray
 - ◆ No merge step needed!
 - $O(n \lg n)$ average case
 - Fast in practice
 - $O(n^2)$ worst case
 - ◆ Naïve implementation: worst case on sorted input
 - ◆ Address this with randomized quicksort

How Fast Can We Sort?

6

- We will provide a lower bound, then beat it
 - by playing a different game
- First, an observation: all of the sorting algorithms so far are *comparison sorts*
 - The only operation used to gain ordering information about a sequence is the **pairwise comparison of two elements**
 - *Theorem*: all comparison sorts are $\Omega(n \log n)$
 - ◆ A comparison sort must do $\Omega(n)$ comparisons (*why?*)
 - ◆ What about the gap between $\Omega(n)$ and $\Omega(n \log n)$

Decision Trees

7

- *Decision trees* provide an abstraction of comparison sorts
 - A decision tree represents the comparisons made by a comparison sort. Everything else is ignored
 - *What do the leaves represent?*
 - leaf is labeled by the permutation of orders that the algorithm determines
 - *How many leaves must there be?*
 - There are $\geq n!$ leaves, because every permutation appears at least once.


Note: Permutations (Appendix C)

8

- A **permutation** of a finite set S is an ordered sequence of all the elements of S , each element appearing exactly once.
-
- For example, if $S = \{a, b, c\}$, then S has 6 permutations: abc , acb , bac , bca , cab , cba
- There are $n!$ permutations of a set of n elements
 - we can choose the first element of the sequence in n ways, the second in $n-1$ ways, the third in $n-2$, etc.

Example: insertion sort for 3 numbers

9



Decision Trees

10

- Decision trees can model comparison sorts. For a given algorithm:
 - One tree for each n
 - Tree paths are all possible execution traces
 - *What's the longest path in a decision tree for insertion sort? For merge sort?*
- *What is the asymptotic height of any decision tree for sorting n elements?*
- Answer: $\Omega(n \log n)$ (proof follows)

Lower Bound For Comparison Sorting

11

- Thm: Any decision tree that sorts n elements has height $\Omega(n \log n)$
- *What's the maximum # of leaves of a binary tree of height h ?*
- Lemma: Any binary tree of height h has $k \leq 2^h$, where k : # of leaves (proof by induction)

Lower Bound For Comparison Sorting

12

- So we have...
 $n! \leq 2^h$
- Taking logarithms:
 $\lg(n!) \leq h$
- Stirling's approximation tells us:

$$n! > \left(\frac{n}{e}\right)^n$$

- Thus:

$$h \geq \log\left(\frac{n}{e}\right)^n$$

Lower Bound For Comparison Sorts

13

- So we have

$$h \geq \log \left(\frac{n}{e} \right)^n$$

$$= n \log n - n \log e$$

$$= \Omega(n \log n)$$

- Thus the minimum height of a decision tree is $\Omega(n \log n)$

Lower Bound For Comparison Sorts

14

- Thus the time to comparison sort n elements is $\Omega(n \log n)$
- Corollary: Heapsort and Mergesort are asymptotically optimal comparison sorts
- “Sorting in linear time” (?)
 - *How can we do better than $\Omega(n \log n)$?*

Sorting In Linear Time

15

- Counting sort
 - No comparisons between elements!
 - *But*...depends on assumption about the numbers being sorted
 - ◆ We assume numbers are in the range $1..k$
 - The algorithm:
 - ◆ Input: $A[1..n]$, where $A[j] \in \{1, 2, 3, \dots, k\}$
 - ◆ Output: $B[1..n]$, sorted (notice: not sorting in place)
 - ◆ Also: Array $C[1..k]$ for auxiliary storage

Counting Sort

16

```
1 CountingSort(A, B, k)
2   for i=1 to k
3       C[i]= 0;
4   for j=1 to n
5       C[A[j]] += 1;
6   for i=2 to k
7       C[i] = C[i] + C[i-1];
8   for j=n downto 1
9       B[C[A[j]]] = A[j];
10      C[A[j]] -= 1;
```


Counting Sort

17

```
1 CountingSort(A, B, k)
2   for i=1 to k
3       C[i] = 0;
4   for j=1 to n
5       C[A[j]] += 1;
6   for i=2 to k
7       C[i] = C[i] + C[i-1];
8   for j=n downto 1
9       B[C[A[j]]] = A[j];
10      C[A[j]] -= 1;
```

Takes time $O(k)$



Takes time $O(n)$



What will be the running time?

Example

18

	1	2	3	4	5	6	7	8
<i>A</i>	2	5	3	0	2	3	0	3

	0	1	2	3	4	5
<i>C</i>	2	0	2	3	0	1

(a)

	0	1	2	3	4	5
<i>C</i>	2	2	4	7	7	8

(b)

	1	2	3	4	5	6	7	8
<i>B</i>							3	

	0	1	2	3	4	5
<i>C</i>	2	2	4	6	7	8

(c)

	1	2	3	4	5	6	7	8
<i>B</i>		0					3	

	0	1	2	3	4	5
<i>C</i>	1	2	4	6	7	8

(d)

	1	2	3	4	5	6	7	8
<i>B</i>		0				3	3	

	0	1	2	3	4	5
<i>C</i>	1	2	4	5	7	8

(e)

	1	2	3	4	5	6	7	8
<i>B</i>	0	0	2	2	3	3	3	5

(f)

Counting Sort

19

- Total time: $O(n + k)$
 - Usually, $k = O(n)$
 - Thus counting sort runs in $O(n)$ time
- But sorting is $\Omega(n \log n)$
 - No contradiction--this is **not a comparison** sort (in fact, there are *no* comparisons at all!)
 - Notice that this algorithm is *stable* (what is that?)
- **Stable:** keys with same value appear in same order in output as they did in input

Counting Sort

20

- Cool! *Why don't we always use counting sort?*
- Because it depends on range k of elements
- *Could we use counting sort to sort 32 bit integers? Why or why not? How many possible (distinct) numbers can we have?*
- Answer: no, k too large ($2^{32} = 4,294,967,296$)

Counting Sort

21

- *How did IBM get rich originally?*
- Answer: punched card readers for census tabulation in early 1900's.
 - In particular, a *card sorter* that could sort cards into different bins
 - ◆ Each column can be punched in 12 places
 - ◆ Decimal digits use 10 places
 - Problem: only one column can be sorted on at a time

Radix Sort

22

- Intuitively, you might sort on the most significant digit, then the second msd, etc.
- Problem: lots of intermediate piles of cards (read: scratch arrays) to keep track of
- Key idea: sort the *least* significant digit first

RadixSort(A, d)

for i=1 to d

StableSort(A) on digit i

Example of Radix Sort

23

329		720		720		329
457		355		329		355
657		436		436		436
839	→	457	→	839	→	457
436		657		355		657
720		329		457		720
355		839		657		839

Radix Sort

24

- *Can we prove it will work?*
- Sketch of an inductive argument (induction on the number of passes):
 - Assume lower-order digits $\{j: j < i\}$ are sorted
 - Show that sorting next digit i leaves array correctly sorted
 - ◆ If two digits at position i are different, ordering numbers by that digit is correct (lower-order digits irrelevant)
 - ◆ If they are the same, numbers are already sorted on the lower-order digits. Since we use a **stable** sort, the numbers stay in the right order

Radix Sort

25

- *What sort will we use to sort on digits?*
- Counting sort is obvious choice:
 - Sort n numbers on digits that range from $1..k$
 - Time: $O(n + k)$
- Each pass over n numbers with d digits takes time $O(n+k)$, so total time $O(dn+dk)$
 - When d is constant and $k=O(n)$, takes $O(n)$ time
- *How many bits in a computer word?*

How to break each key into digits?

26

- n words
- b bits/word
- Break into r -bit digits. Have $d = \lceil b/r \rceil$ digits
- Use counting sort with $k = 2^r - 1$
- Example: 32-bit words, 8-bit digits. $b = 32$, $r = 8$, $d = 32/8 = 4$, $k = 2^8 - 1 = 255$.
- • Time = $\Theta(b/r(n + 2^r))$.
- Choose $r \approx \log n$ gives: $\Theta(bn / \log n)$.

Bucket Sort

27

- Assumes the input is generated by a random process that distributes elements uniformly over $[0, 1)$.

Idea:

- Divide $[0, 1)$ into n equal-sized *buckets*.
- Distribute the n input values into the buckets.
- Sort each bucket.
- Then go through buckets in order, listing elements in each one.

Input: $A[1 \dots n]$, where $0 \leq A[i] < 1$ for all i .

Auxiliary array: $B[0 \dots n - 1]$ of linked lists, each list initially empty.

Bucket sort Code

28


```
BUCKET-SORT(A, n)
for i  $\leftarrow$  1 to n
    insert A[i] into list B[ $\lfloor n \cdot A[i] \rfloor$ ]
for i  $\leftarrow$  0 to n - 1
    sort list B[i] with insertion sort
    concatenate lists B[0], B[1], . . . , B[n
- 1] together in order
return the concatenated lists
```

Bucket Sort Example

(29)

<i>A</i>	
1	.78
2	.17
3	.39
4	.26
5	.72
6	.94
7	.21
8	.12
9	.23
10	.68

(a)



(b)

Correctness

30

- Consider $A[i]$, $A[j]$.
- Assume without loss of generality that $A[i] \leq A[j]$.
Then $\lfloor n \cdot A[i] \rfloor \leq \lfloor n \cdot A[j] \rfloor$.
- Case1: $A[i]$ is placed into the same bucket as $A[j]$ or
- Case2: into a bucket with a lower index (Case 2).

If same bucket (Case1), insertion sort fixes up.

If earlier bucket (Case2), concatenation of lists fixes up.

Analysis:

31

- Relies on *no bucket* getting too many values.
- All lines of algorithm except insertion sorting take $\Theta(n)$ altogether.
- Intuitively, if each bucket gets a constant number of elements, it takes $O(1)$ time to sort each bucket $\Rightarrow O(n)$ sort time for all buckets.
- We “expect” each bucket to have few elements, since the average is 1 element per bucket.
- Uniform input distribution has $O(1)$ bucket size and expected time is $O(n)$
- Later in Hash Tables again the same idea

Structures...

32

- Done with sorting and order statistics
- Next part is *data structures Ch 10 (skim)*
- *Ch 11*